

OSL Practice Problem Statement AY 2025-26

1.1 Write a shell script to implement an address book (address.txt) that contains ID, Name, and Phone with the following functions:

1. Search Address Book
2. Add an address book entry
3. Remove an address book entry
4. Quit the program.

Suggested record format: Use semi-colon (;) to separate fields.

1.2 Write a script program to create a phonebook. Using the phonebook, perform the following tasks:

1. Add a new name to a phone book file
2. Display all matches to a name or phone number
3. Sort the phone book by the last name
4. Delete an entry

Suggested record format: Use tabs to separate fields.

1.3 Create an address book program using the bourne-again shell. It should use functions to perform the required tasks. It should be menu-based, allowing you the options of:

1. Search address book
2. Add entries
3. Remove / edit entries.

You will also need a "display" function to display a record or records when selected.

Suggested record format: Use colons to separate fields.

1.4 Create a file management program using the bourne-again shell. It should use functions to perform the required tasks. It should be menu-based, allowing you the options of:

1. Test if file exists
2. Read a file
3. Delete a file
4. Display a list of files

e.g. book.txt and test.txt files are available in present working directory. The filename = "book.txt" contents:

1. Pro AngularJS
2. Learning JQuery
3. PHP Programming

1.5 Write a shell script

- a) that accepts a file name, starting and ending line numbers as arguments and displays all the lines between the given line numbers.
- b) that deletes all lines containing a specified word in one or more files supplied as arguments to it.
- c) that takes two numbers as arguments and output their sum using i) bc ii) expr. Include error checking to test if two arguments were entered.

1.6 Write a shell script

- a) that uses find to look for a file and echo a suitable msg if the file is not found. You must not store the find output in a file.
- b) that takes a command-line argument and reports on whether it is directory, a file, or something else.
- c) that accepts one or more file name as arguments and converts all of them to uppercase, provided they exist in the current directory.

2.1 Implement the C program to accept 'n' integers to be sorted. Main function creates child process using fork system call. Parent process sorts the integers using **bubble sort** and waits for child process using wait system call. Child process sorts the integers using **insertion sort**. Also demonstrate zombie and orphan states.

2.2 Implement the C program in which main program accepts the integers to be sorted. Main program uses the fork system call to create a new process called a child process. Parent process sorts the integers using **merge sort** and waits for child process using wait system call to sort the integers using **quick sort**. Also demonstrate zombie and orphan states.

2.3 Implement the C program in which main program accepts an integer array. Main program uses the fork system call to create a new process called a child process. Parent process sorts an integer array and passes the sorted array to child process through the command line arguments of execve system call. The child process uses execve system call to load new program that uses this sorted array for performing the binary search to search the item in the array.

2.4 Write a C program using the fork() system call that generates the sequence in the child process. The starting number will be provided from the command line. For example, if 8 is passed as a parameter on the command line, the child process will output 8, 4, 2, 1. i.e., when we take any positive integer n and apply the following algorithm, it will eventually reach 1:

$$n = \begin{cases} n/2, & \text{if } n \text{ is even} \\ 3 \times n + 1, & \text{if } n \text{ is odd} \end{cases}$$

Because the parent and child processes have their own copies of the data, it will be necessary for the child to output the sequence. Have the parent invoke the wait() call to wait for the child process to complete before exiting the program. Perform necessary error checking to ensure that a positive integer is passed on the command line.

2.5 Write a C program using the fork() system call that generates the prime number sequence in the child process. The number will be provided from the command line. For example, if 10 is passed as a parameter on the command line, the child process will output 2, 3, 5, 7, 11, 13, 17, 19, 23, 29. Because the parent and child processes have their own copies of the data, it will be necessary for the child to output the sequence. Have the parent invoke the wait() call to wait for the child process to complete before exiting the program. Perform necessary error checking to ensure that a positive integer is passed on the command line.

3.1 Write a menu-driven C program to simulate the following CPU scheduling algorithms to find average turnaround time and average waiting time.

a) FCFS b) SJF (non-preemptive)

Expected Output (e.g):

Enter the number of processes:3

Process	Arrival time	Burst time
p1	0	7
p2	2	4
p3	4	1
p4	5	4

Enter the choice: 1. FCFS 2. SJF (non-preemptive) 3. Exit

Average waiting time for <FCFS/SJF> is:

Average turn-around time for <FCFS/SJF> is:

GANTT CHART (e.g):

	p3		p2		p1	
0	5	9	12			

3.2 Write a menu-driven C program to simulate the following CPU scheduling algorithms to find average turnaround time and average waiting time.

a) FCFS b) SJF (preemptive)

Expected Output (e.g):

Enter the number of processes:3

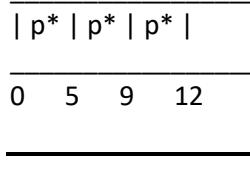
Process	Arrival time	Burst time
p1	0	7
p2	2	4
p3	4	1
p4	5	4

Enter the choice: 1. FCFS 2. SJF (preemptive) 3. Exit

Average waiting time for <FCFS/SJF> is:

Average turn-around time for <FCFS/SJF> is:

GANTT CHART:



3.3 Write a menu-driven C program to simulate the following CPU scheduling algorithms to find average turnaround time and average waiting time.

a) FCFS b) Round-robin (preemptive) with time quantum = 2 sec.

Expected Output (e.g):

Enter the number of processes:3

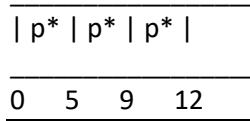
Process	Arrival time	Burst time
p1	0	7
p2	2	4
p3	4	1
p4	5	4

Enter the choice: 1. FCFS 2. Round-robin (preemptive) 3. Exit

Average waiting time for <FCFS/RR> is:

Average turn-around time for <FCFS/RR> is:

GANTT CHART:



4.1 The problem below describes two processes, the producer and the consumer, who share a common, fixed-size buffer used as a queue.

The producer is Mr. Simpson whose job is to bake Pizza and consumer is Joey Tribbiani who consumes Pizza at the same time. Both of them share common Pizza counter to interact with each other. The problem is to make sure that the Simpson won't try to add more pizza on Pizza-counter if it's full. He needs to wait until Joey consumes Pizza. Similarly, Joey can't consume pizza from an empty counter. He needs to wait until Mr. Simpson adds Pizza on counter.

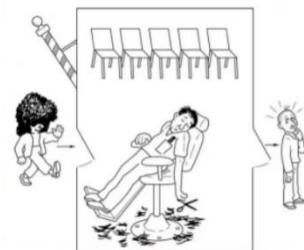
Implement a solution in C using POSIX threads that coordinates the activities of the producer Simpson and the consumer Joey with counting semaphores and mutex.

4.2 A university computer science department has a teaching assistant (TA) who helps undergraduate students with their programming assignments during regular office hours. The TA's office is rather small and has room for only one desk with a chair and computer. There are three chairs in the hallway outside the office where students can sit and wait if the TA is currently helping another student. When there are no students who need help during office hours, the TA sits at the desk and takes a nap. If a student arrives during office hours and finds the TA sleeping, the student must awaken the TA to ask for help. If a student arrives and finds the TA currently helping another student, the student sits on one of the chairs in the hallway and waits. If no chairs are available, the student will come back at a later time. Using POSIX threads, mutex locks, and semaphores, implement a solution that coordinates the activities of the TA and the students.

4.3 The Producer generates an integer between 0 and 9 (inclusive), stores it in a file. To make the synchronization problem more interesting, the Producer sleeps for a random amount of time between 0 and 100 milliseconds before repeating the number-generating cycle. The Consumer consumes all integers from the same file (the exact same file into which the Producer put the integers in the first place) as quickly as they become available. The activities of the Producer and the Consumer must be synchronized in a following way: The two threads must do some simple coordination. That is, the Producer must have a way to indicate to the Consumer that the value is ready, and the Consumer must have a way to indicate that the value has been retrieved. Implement the solution using mutex, and binary semaphore — to help threads wait for a condition and notify other threads when that condition changes.

4.4 Sleeping Barber Problem:

Suppose a customer arrives and notices that the barber is busy cutting the hair of another customer, so he goes to the waiting room. While he is on his way to the waiting room, the barber finishes his job and sees the waiting room for other customers. But he(the barber) finds no one in the waiting room (as the customer has yet not arrived in the waiting room), so he sits down in his chair(barber chair) and sleeps. Now the barber is waiting for new customers to wake him up, and the customer is waiting as he thinks the barber is busy.



Implement the solution, using three semaphores, one for customers(for counts of waiting for customers), one for barber(a binary semaphore denoting the state of the barber, i.e., 0 for idle and 1 for busy), and a mutual exclusion semaphore, mutex for seats.

4.5 A database must be shared by numerous concurrent processes, some of which may simply want to read the database, while others may wish to update (read and write) the database. We differentiate between these two processes by referring to the former as Readers and the latter as Writers. There will be no negative consequences if two readers access the shared data simultaneously. If a writer and another thread (either a reader or a writer) access the common data simultaneously, chaos may follow.

To avoid these problems, implement a solution that the writers have exclusive access to the shared database by using mutex, and semaphore with readers being prioritized.

4.6 A real-world example of the readers-writers problem is an airline reservation system:

- Readers: want to read flight information
- Writers: want to make flight reservations
- Potential problem: if readers and writers can access the shared data simultaneously then readers/writers may view flights as being available when they've actually just been booked.

Implement the solution using mutex and semaphore that controls access to the reader count and database given the priority to readers over writers.

4.7 Consider an airline reservation system in which many clients are attempting to book seats on particular flights between particular cities. All of the information about flights and seats is stored in a common database in memory. The database consists of many entries, each representing a seat on a particular flight for a particular day between particular cities. In a typical airline reservation scenario, the client will probe around in the database looking for the "optimal" flight to meet that client's needs. So a client may probe the database many times before deciding to try and book a particular flight. A seat that was available during this probing phase could easily be booked by someone else before the client has a chance to book it after deciding on it. In that case, when the client attempts to make the reservation, the client will discover that the data has changed and the flight is no longer available. The client probing around the database is called a reader. The client attempting to book the flight is called a writer. Clearly, any number of readers can be probing shared data at once, but each writer needs exclusive access to the shared data to prevent the data from being corrupted.

Implement a C program that launches multiple reader threads and multiple writer threads, each attempting to access a single reservation record. Implement a version of your program that uses mutex and counting semaphore to enforce a disciplined protocol for readers and writers accessing the shared reservation data. In particular, your program should favor readers to access the shared data simultaneously over writers.

4.8 Assume that a finite number of resources of a single resource type must be managed. Processes may ask for a number of these resources and will return them once finished. As an example, many commercial software packages provide a given number of licenses, indicating the number of applications that may run concurrently. When the application is started, the license count is decremented. When the application is terminated, the license count is incremented. If all licenses are in use, requests to start the application are denied. Such requests will only be granted when an existing license holder terminates the application, and a license is returned. Using a semaphore or mutex lock, fix the race condition. The calling process may be blocked until sufficient resources are available. When a process wishes to obtain a number of resources, it invokes the decrease count() function: decrease available resources by count resources and return 0 if sufficient resources available.

When a process wants to return a number of resources, it calls the increase count() function: increase available resources by count

5.1 Develop a C program to implement Banker's algorithm. Assume suitable input required to demonstrate the results. Using the banker's algorithm, determine whether or not the state is unsafe. If the state is safe, illustrate the order in which the processes may complete. Otherwise, illustrate why the state is unsafe.

5.2 Develop a C program to simulate Banker's Algorithm for Deadlock Avoidance with following requirements:

1. Demonstrate safe and unsafe state of the system?
2. Demonstrate Grant and non-grant of new resource request.

Note: Your program may read the current state of the system from a file called "state.txt". In this file, number of processes, number of resources, Allocation Matrix, Max and Available Resources Matrix may be presented in the order below:

State.txt contents:-->

Number of Processes	5
Number of Resources	4

```
3 2 0 0
0 1 1 2
2 1 0 0
0 0 1 0
2 1 1 1
```

```
3 0 1 1
0 1 0 0
1 1 1 0
1 1 0 1
0 0 0 0
```

```
1 0 2 1
```

6.1 Write a C program to simulate page replacement algorithms a) FIFO b) LRU

First, generate a random page-reference string where page numbers range from 0 to 9. Apply the random page-reference string to each algorithm and record the number of page faults incurred by each algorithm. Implement the replacement algorithms so that the number of page frames can vary from 1 to 7.

6.2 Consider the page reference string of size 12: 1, 2, 3, 4, 1, 2, 5, 1, 1, 2, 3, 4, 5 with frame size 3 and 4 respectively. Write a C program to simulate page replacement algorithms a) Optimal b) LRU.

7.1 Design a program using ordinary pipes in which one process sends a string message to a second process, and the second process reverses the case of each character in the message and sends it back to the first process. For example, if the first process sends the message Hi There, the second process will return hI tHERE. This will require using two pipes, one for sending the original message from the first to the second process and the other for sending the modified message from the second to the first process. You can write this program using LINUX or Windows pipes.

7.2 Write a C program in which the parent process accepts a string from the user and writes it into a pipe. The child process reads the string, reverses it, and displays the reversed output on the console. Ensure proper closing of pipe ends and error handling during communication.

7.3 Implement two independent processes, Sender and Receiver, using Message Queues (FIFO). The Sender accepts messages from the user and sends them to the queue with different message types. The Receiver reads each message, converts it to uppercase, and displays it on the screen. Ensure proper cleanup of message queues after execution.

7.4 Develop Client and Server processes using Shared Memory.

The Server creates a shared memory segment and writes a message to it. The Client reads the message from the shared memory segment and displays it on the screen. Include appropriate synchronization (using semaphores or sleep/wait) to avoid data inconsistency.

8.1 Implement and compare the total head movement for a given sequence of disk requests using two specific scheduling algorithms:

- System Setup: A disk with cylinders numbered 0 to 499.
- Current Head Position: 85
- Pending Requests (FIFO order): 10, 229, 39, 400, 18, 145, 120, 480, 20, 250
- Algorithms to Implement:
 - C-SCAN (Assume initial movement is towards 499).
 - C-LOOK (Assume initial movement is towards 499).
- Output: Calculate and display the sequence of cylinder movements and the Average Seek Distance for both C-SCAN and C-LOOK. Conclude which algorithm performed better for this request set.

8.2 A disk drive has 5000 cylinders, numbered 0 to 4999. The drive is currently serving a request at cylinder 143. The previous request was at cylinder 125. The queue of pending requests, in FIFO order, is:

86, 1470, 913, 1774, 948, 1509, 1022, 1750, 130

Calculate the average seek length (in cylinders) that the disk arm moves to satisfy all of the pending requests for each of the following disk-scheduling algorithms.

SCAN (Assume the head is currently moving **toward 4999**)

LOOK (Assume the head is currently moving **toward 4999**)

8.3 A disk drive has 5000 cylinders, numbered 0 to 4999. The drive is currently serving a request at cylinder 143. The previous request was at cylinder 125. The queue of pending requests, in FIFO order, is:

86, 1470, 913, 1774, 948, 1509, 1022, 1750, 130

Calculate the average seek length (in cylinders) that the disk arm moves to satisfy all of the pending requests for each of the following disk-scheduling algorithms:

- **C-SCAN** (Assume the head is currently moving **toward 4999**)
- **C-LOOK** (Assume the head is currently moving **toward 4999**)

General Oral Questions

- What is the race condition?
- What is a binary semaphore?
- What is the difference between Semaphore and Mutex.
- Explain the concept of semaphore?
- Explain wait and signal functions associated with semaphores.
- Where the binary and counting semaphores are used in the solution of classical problems.
- Is dining philosopher problem represents processes synchronization or deadlock situation?
- How to avoid deadlock in dining philosopher problem.
- How many operations Banker's algorithm required to perform in worst case scenario?
- What is the cause of thrashing?
- Why LRU and Optimal Page Replacement algorithm doesn't suffer from Belady's anomaly.
- What is the fundamental difference between an unnamed pipe and a FIFO (named pipe) in Linux, and when would you choose one over the other?
- What is the main advantage of using Shared Memory over message passing mechanisms like pipes or message queues? What is its major drawback?
- Why is a separate synchronization mechanism (like a semaphore or mutex) absolutely essential when using shared memory?
- Define a **race condition** and provide a simple, real-world example different from the bank account simulation.
- What is the key difference between a Mutex and a counting Semaphore? Which one is better suited for protecting the bank balance variable?
- Name the four necessary and sufficient conditions for a deadlock to occur.
- Briefly distinguish between deadlock and livelock.
- Why is the Banker's Algorithm primarily a theoretical tool rather than a widely implemented feature in general-purpose operating systems?
- What are the two main performance metrics that disk scheduling algorithms aim to optimize?
- Which of the six common disk scheduling algorithms (FIFO, SSTF, SCAN, C-SCAN, LOOK, C-LOOK) is most prone to causing starvation for requests located at the edges of the disk? Explain why.
- Why does C-SCAN generally provide a more uniform waiting time than SCAN, despite often resulting in a higher total seek time?
- Based on our earlier comparison, if disk load is light and requests are clustered, why is LOOK preferred over SCAN?